

Assignment 4: Shallow Waters

Practical information

Deadline: Monday 3/11 23.59

Resources:

- ERDA for file storage
- Nvidia profiler to determine the parallelisation bottlenecks
- Jupyter for the Terminal to access DAG Nvidia vGPU instances

Handin:

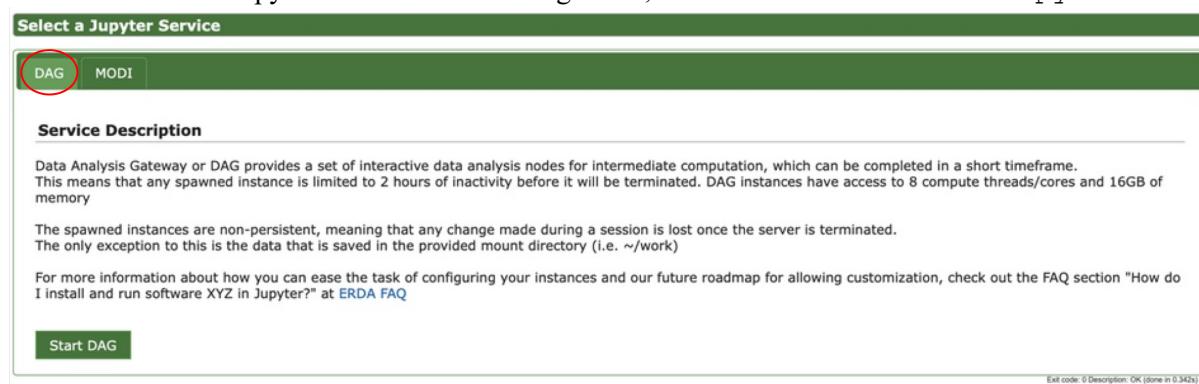
- Total assignment: a report of up to 3 pages in length (excluding the code)
- Use the template on Absalon to include your code in the report

Introduction

The Shallow Water (SW) model (see PHPC section 13.3 or the lecture notes on Absalon; first 3 pages are enough) is the simplest numerical representation of the ocean. Still, it has reasonable precision when used to predict the evolution of storm surges or Tsunamis. Moreover, it illustrates nicely the functioning and parallelization of stencil operations.

DAG

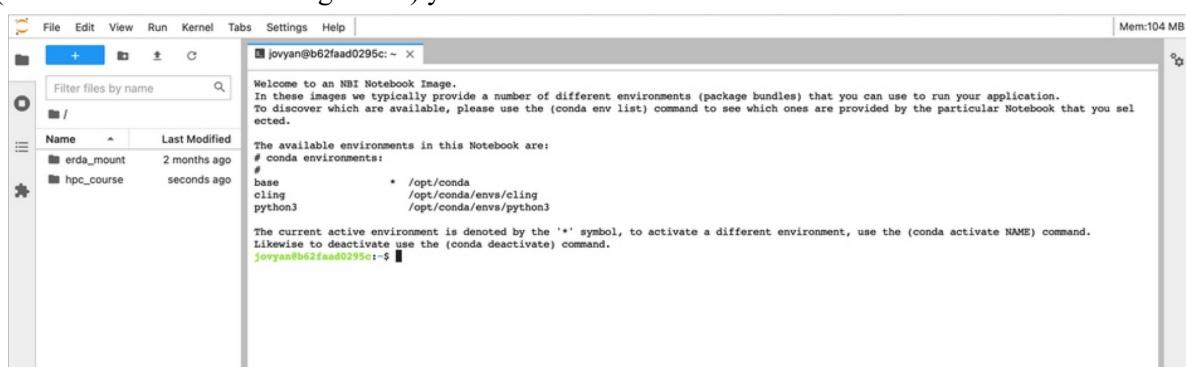
For this assignment we need nvc++ to compile and DAG for source code profiling and running the benchmarks. For the python version of the assignment, we need to do “conda install cupy”.



The screenshot shows a web-based interface for selecting a Jupyter service. At the top, there is a green header bar with the text "Select a Jupyter Service". Below this, there are two tabs: "DAG" and "MODI". The "DAG" tab is highlighted with a red circle around it. The main content area is titled "Service Description". It contains the following text:
Data Analysis Gateway or DAG provides a set of interactive data analysis nodes for intermediate computation, which can be completed in a short timeframe. This means that any spawned instance is limited to 2 hours of inactivity before it will be terminated. DAG instances have access to 8 compute threads/cores and 16GB of memory.
The spawned instances are non-persistent, meaning that any change made during a session is lost once the server is terminated. The only exception to this is the data that is saved in the provided mount directory (i.e. ~/work).
For more information about how you can ease the task of configuring your instances and our future roadmap for allowing customization, check out the FAQ section "How do I install and run software XYZ in Jupyter?" at [ERDA FAQ](#).

At the bottom left, there is a green button labeled "Start DAG". At the bottom right, there is a small text "Exit code: 0 Description: OK (done in 0.342s)".

You can read more about DAG in the user guide: <https://erda.dk/public/ucph-erda-user-guide.pdf>
Spin up a Jupyter session on DAG selecting the “HPC GPU notebook” notebook image. In the terminal (or the folder view on the right side) you can see a number of folders.



The different folders contain:
erda_mount: your own files.
hpc_course: not relevant.

Preparations

Start by updating the exercise folder in your storage area and enter in to the folder. You can write ‘ls’ to get a file listing of the folder. Steps below are for C++ and Python. Fortran is very similar to C++.

```
cd erda_mount/ahpc  
git pull  
cd week4/cpp  
cd week4/fortran  
or  
cd week4/python
```

You can write ‘ls’ to get a file listing of the folder. For C++ it looks like

```
Makefile (only for C++ and Fortran)  
sw_parallel.cpp/f90/py  
sw_sequential.cpp/f90/py  
visualize.ipynb  
run_sw.sh
```

For C++ and Fortran, before you can run the code, you need to compile it. This can be done by running `make` in the terminal. The `sw_sequential.cpp/f90/py` code is identical to `sw_parallel.cpp/f90/py` and is there (with produced corresponding binaries `sw_sequential` and `sw_parallel`) to give you a backup. The `visualize.ipynb` is for Shallow Waters model output visualisation and analysis. The `run_sw.sh` script can help you with launching the code on a restricted number of streaming multiprocessors to do weak scaling.

To run the code for 500 time-steps on the CPU on DAG and write the model output in ASCII file to your storage you can do:

```
./sw_sequential --iter 500 --out sw_output.txt
```

Or in Python

```
python ./sw_sequential --iter 500 --out sw_output.txt
```

To run the code for 500 time-steps on the GPU and write the model output in ASCII file to your storage you can do:

```
./sw_parallel --iter 500 --out sw_output_gpu.txt
```

Or in Python

```
python ./sw_parallel --iter 500 --out sw_output_gpu.txt
```

You should at least once try to visualize your output. Visualizing your output is also good for validation.

Nvidia profiler

NVIDIA profiler (nsys) enables you to understand and optimize the performance of your GPU application. An example of command-line nsys profiler output for parallelised SW model using CuPy and profiler decorators is given below. In the different sections there are:

- [3/8]: Time spent, number of calls etc in the decorated functions. Directly related to source code
- [4/8]: Timing for OS and system calls, can normally be ignored
- [5/8]: Timing for CUDA specific functions
- [6/8]: Timing for individual kernels. Generated by CuPy / OpenACC
- [7/8]: Memory transfer statistics from host to device and device to host
- [8/8]: GPU memory usage by the program

All the information is quite useful for understanding what contributes to runtime, and can be sueful to e.g. detect transfer of small data done implicitly by either CuPy or OpenACC.

```
% nsys profile --stats=true /lustre/astro/troels/teaching/ahpc/venv/ahpc/bin/python sw_parallel_solution.py
checksum: 16470.993291673218
elapsed time: 0.19375168485566979 sec
Generating '/tmp/nsys-report-b74e.qdstrm'
[1/8] [=====100%] report3.nsys-report
[2/8] [=====100%] report3.sqlite
[3/8] Executing 'nvtx_sum' stats report

Time (%) Total Time (ns) Instances Avg (ns) Med (ns) Min (ns) Max (ns) StdDev (ns) Style Range
----- -----
 62.5    46103239      1000   46103.2   40831.0    39018   3556885   112000.1 PushPop integrate
19.0     14024712      2000   7012.4    6109.0    5629   930805   21758.4 PushPop exchange_horizontal_ghost_lines
18.4     13606419      2000   6803.2    6079.0    5628   852759   20006.1 PushPop exchange_vertical_ghost_lines
 0.1      74390         1    74390.0   74390.0   74390   74390       0.0 PushPop CCCL:cub::DeviceReduce::Sum

[4/8] Executing 'osrt_sum' stats report

Time (%) Total Time (ns) Num Calls Avg (ns) Med (ns) Min (ns) Max (ns) StdDev (ns) Name
----- -----
 38.5    566105651      15  37740376.7  9627771.0    2143   164757075  52899004.9 poll
.....output related to OS calls and libraries.....

[5/8] Executing 'cuda_api_sum' stats report

Time (%) Total Time (ns) Num Calls Avg (ns) Med (ns) Min (ns) Max (ns) StdDev (ns) Name
----- -----
 50.3    146878293      12  12239857.8  15965730.5   34551   16220945   6820732.2 cudaMemcpyAsync
44.2     129041961      12  10753496.8   98942.0    5869   128119089   36960595.4 cudaMalloc
 4.8      13897744      6027   2305.9    2233.0    2053   19379       656.7 cuLaunchKernel
 0.4      1275536       14   91109.7   78567.0    75352   247358   45061.5 cuModuleLoadData
 0.1      400376       9    44486.2   42744.0   31397   82612   16256.5 cuModuleUnload
 0.0      122901       440    279.3    230.0      60    4126       275.3 cuGetProcAddress_v2
 0.0      57285        2    28642.5   28642.5    25057   32228   5070.7 cuLaunchKernel
 0.0      45998        3    15332.7   7602.0    6499   31897   14355.7 cudaMemcpyAsync
 0.0      45298        1    45298.0   45298.0   45298   45298       0.0 cudaMemcpyGetInfo
 0.0      39958        12   3329.8    3204.5    2463   5168       771.3 cudaMemcpySynchronize
 0.0      13190       12    1099.2   936.0    531    2995       671.5 cudaMemcpyIsCapturing_v10000
 0.0      6260         2    3130.0   3130.0    160    6100   4200.2 cuModuleGetLoadingMode
 0.0      5278         2    2639.0   2639.0    2574   2704       91.9 cuInit
 0.0      721          1    721.0    721.0    721    721       0.0 cuLibraryGetKernel
 0.0      321          2    160.5   160.5    160    161       0.7 cuKernelGetName

[6/8] Executing 'cuda_gpu_kern_sum' stats report

Time (%) Total Time (ns) Instances Avg (ns) Med (ns) Min (ns) Max (ns) StdDev (ns) Name
----- -----
 49.1    89292515      1000   89292.5   89824.0   84192   90240   1622.1 velocity_update
46.3     84264150      1000   84264.2   84768.0   79520   85153   1537.5 elevation_update
 2.2      3995188      2000   1997.6   1984.0    1791   2272       104.7 exchange_horizontal_ghost_lines
 2.1      3824868      2000   1912.4   1920.0    1759   2144       44.4 exchange_vertical_ghost_lines
 0.2      330847       14   23631.9   25168.0   12672   26432   4525.3 cupy_copy_float64_float64
 0.0      65345        1    65345.0   65345.0   65345   65345       0.0 cupy_exp_float64_float64
 0.0      53472        2    26736.0   26736.0   26528   26944   294.2 cupy_multiply_float64_float64_float64
 0.0      38880        1    38880.0   38880.0   38880   38880       0.0 cupy_add_float64_float64_float64
 0.0      29119        3    9706.3   1408.0    1375   26336   14401.7 cupy_multiply_float_float64_float64
 0.0      16032        1    16032.0   16032.0   16032   16032       0.0 void
cub::CUB_200800_SM_750_800_860_890_900_1000_1200::DeviceReduceKernel<cuib::CUB_200800_SM_750_80...
 0.0      3744         2    1872.0   1872.0    1856   1888       22.6 cupy_true_divide_float64_float64
 0.0      3040         2    1520.0   1520.0    1344   1696       248.9 cupy_subtract_float64_float64
 0.0      2944         2    1472.0   1472.0    1248   1696       316.8 cupy_arange_float_float64
 0.0      2432         1    2432.0   2432.0    2432   2432       0.0 void
cub::CUB_200800_SM_750_800_860_890_900_1000_1200::DeviceReduceSingleTileKernel<cuib::CUB_200800...

[7/8] Executing 'cuda_gpu_mem_time_sum' stats report

Time (%) Total Time (ns) Count Avg (ns) Med (ns) Min (ns) Max (ns) StdDev (ns) Operation
----- -----
 99.6    9144887      12  762073.9   843314.0   1440   863362   241257.8 [CUDA memcpy Device-to-Host]
 0.4      38624        3  12874.7   13248.0   11936   13440   818.6 [CUDA memset]

[8/8] Executing 'cuda_gpu_mem_size_sum' stats report

Total (MB) Count Avg (MB) Med (MB) Min (MB) Max (MB) StdDev (MB) Operation
----- -----
 92.275     12    7.690   8.389   0.000   8.389   2.422 [CUDA memcpy Device-to-Host]
 25.166      3    8.389   8.389   8.389   8.389   0.000 [CUDA memset]
```

Task 1: OpenACC / CuPy parallelise the program

The key challenge is to identify which parts of the code can reasonably be executed by the GPUs. For C++ and Fortran you should find suitable OpenACC directives and clauses for optimal parallelization. For Python you should use CuPy instead of NumPy. Use the profiler to determine the bottlenecks. For C++/Fortran, once the code is working, store it, and play around a bit with the #pragma and see if you can improve on your first try. For Python / CuPy once the code has been moved to the GPU consider if you can use fusing of operations or kernels to reduce the number of GPU invocations. In both cases streams may be useful for outputting data to I/O. To complete this task you should attach your code, and present your strategy for GPU parallelising the code. Add central parts of the nsys profiler output to show how your code performs and discuss the results.

Task 2: Weak scaling and asymptotic performance

Measure the weak scaling of your program using the `run_sw.sh` script. The script allows you to constrain the number of streaming multi-processors to an even number (from 2 to 14) using the CUDA Multi-Process Service daemon. Change the variables `NX` and `NY` to measure the weak scaling, you should think about and explain your choice of number of grid cells and how this map to the available vGPUs. Some key figures to note: You have 14 streaming multiprocessors (compute units or SMs) available, the maximum number of threads per thread-block is 1024 and each multiprocessor can handle at most 2048 threads simultaneously. Discuss briefly your figure.

What we often do when running on GPUs is to figure out what is the minimum size of the workload to make good use of the GPU. This can be determined by running on the full GPU while increasing the workload until the run-time increases linearly with the size of the workload. Make a plot of the asymptotic performance measured as the number of nano-seconds it takes to update a single grid-cell as a function of the number of grid cells. Based on your figure, what is a reasonable grid-size to run with on the ERDA GPU to make good use of it?